



# CLARION: Sound and Clear Provenance Tracking for Microservice Deployments

XUTONG CHEN<sup>1</sup>, HASSAAN IRSHAD<sup>2</sup>,  
YAN CHEN<sup>1</sup>, ASHISH GEHANI<sup>2</sup>,  
VINOD YEGNESWARAN<sup>2</sup>

NORTHWESTERN UNIVERSITY<sup>1</sup>,  
SRI INTERNATIONAL<sup>2</sup>

# Microservices, Containers and Provenance Tracking

## Microservices

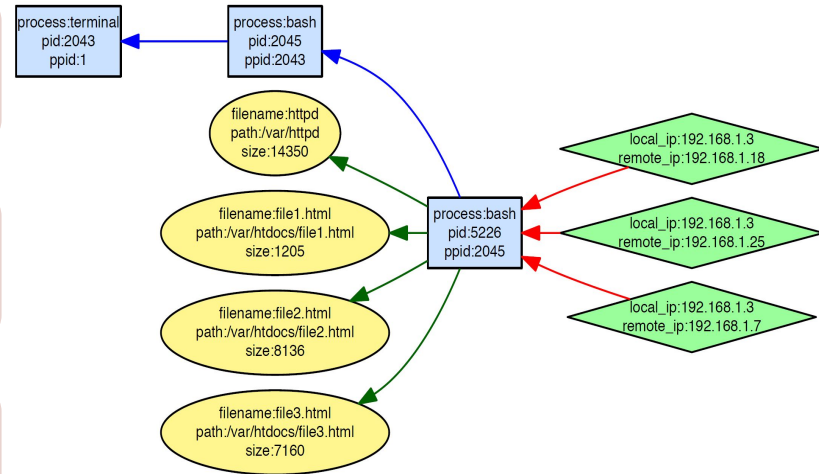
- Enable better resource utilization
- Increasing practical adoption

## Containers

- Portable and lightweight software isolation technique
- New security challenges

## Provenance Tracking


- State-of-the-art host forensic monitoring solution
- Transferred from traditional OS scenario to container scenario



A sample provenance graph generated by SPADE, Middleware'12

# Motivating Example: Insider Attack

---



- Three major steps of a trivial insider attack is shown on the left.
- Perform all three steps in both the host and the container.
- We drill down just on the first step in the next slide for simplicity.
- Motivate our work by illustrating limitations in the provenance graphs from three contemporary solutions.

# Existing Provenance Tracking Solutions

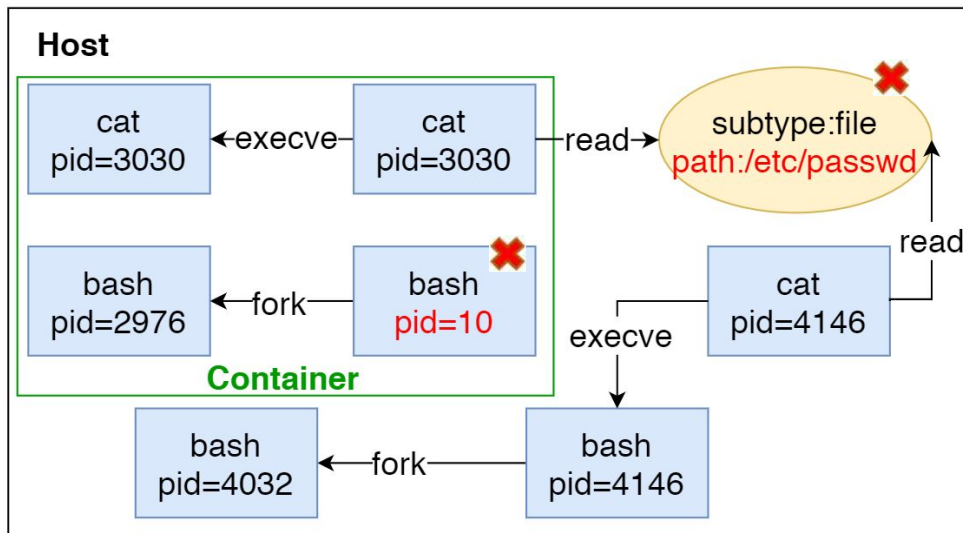
**Namespace-unaware**

**Container-aware**

*Winnower: NDSS 2018*

Fail on soundness

- Fragmentation in *bash-cat* provenance
- Ambiguities on */etc/passwd*



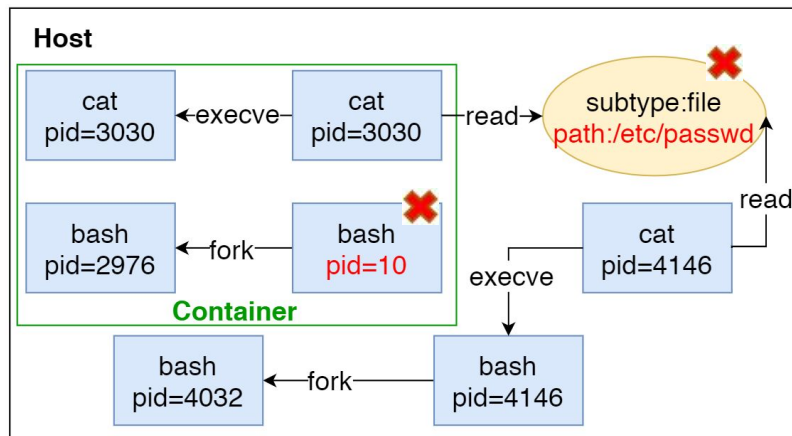
# Existing Provenance Tracking Solutions

## Namespace-unaware Container-aware

Winnower: NDSS 2018

Fail on soundness

- Fragmentation in *bash-cat* provenance
- Ambiguities on */etc/passwd*

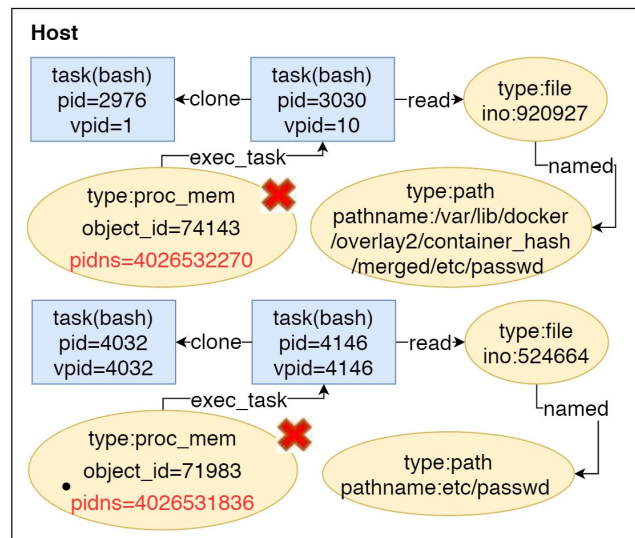


## Namespace-aware Container-unaware

Camflow: SoCC 2017

Fail on clarity

- Unclear insight about the container



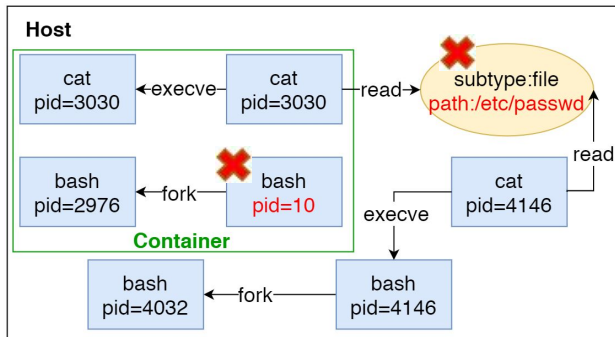
# Existing Provenance Tracking Solutions

## Namespace-unaware Container-aware

*Winnower: NDSS 2018*

Fail on soundness

- Fragmentation in *bash-cat* provenance
- Ambiguities on */etc/passwd*

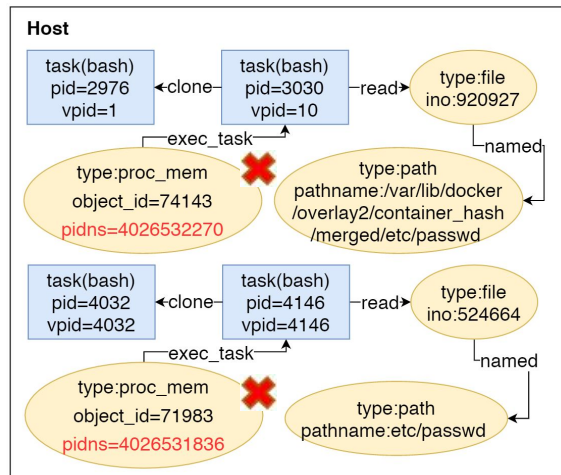


## Namespace-aware Container-unaware

*Camflow: SoCC 2017*

Fail on clarity

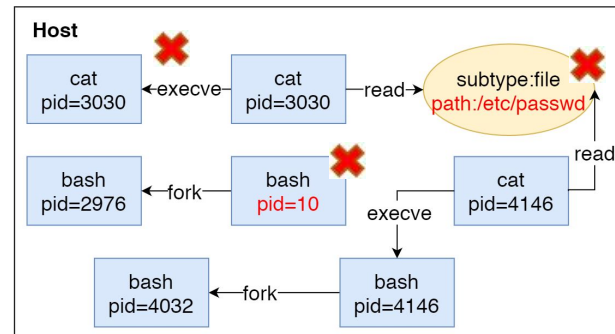
- Unclear insight about the container



## Namespace-unaware Container-unaware

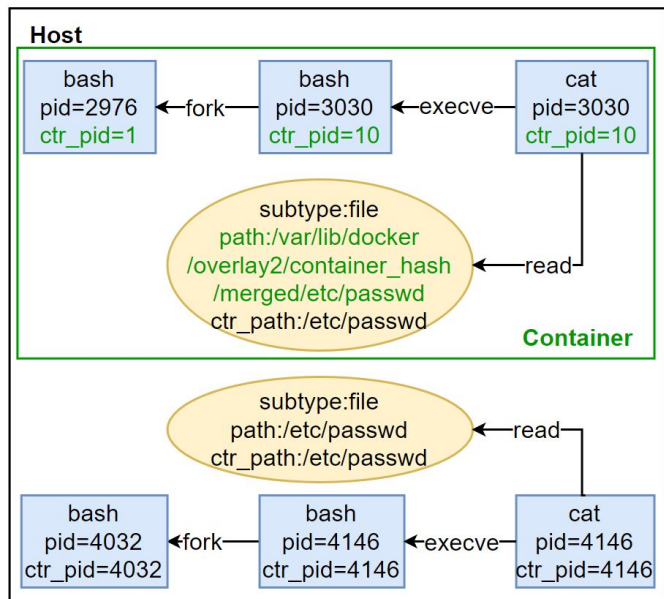
*SPADE: Middleware 2012*

Fail on both soundness  
and clarity



# CLARION Solution and Key Contributions

## Namespace-aware Container-aware



- The first in-depth analysis of the implications of namespaces on provenance tracking.
- CLARION: A **namespace-aware** and **container-aware** provenance tracking solution.
  - Extension of the SPADE Provenance Engine to address clarity and soundness issues.
  - Linux kernel module, netfilter hooks, and modifications to SPADE application-level handlers.
- Comprehensive evaluation of the effectiveness, efficiency, and generality.

# Soundness Challenge: Inconsistency in low-level events

---

- *Inconsistency in low-level events* occur when low-level events report data values from varying (host / container) contexts that lead to vertex splitting in provenance graphs.
- Consider the “clone” Linux audit event as an example, its “pid” field (5903) is in host context but its “exit” field (2), which is also a PID, is affected by pid namespace and in container context.
- If provenance tracking system directly uses those two data fields to generate provenance, *graph fragmentation* will occur.

```
type=SYSCALL msg=audit(1567029444.851:431219): arch=c000003e syscall=56  
success=yes exit=2 a0=3d0f00 a1=7f81aa8f8fb0 a2=7f81aa8f99d0  
a3=7f81aa8f99d0 items=0 ppid=5880 pid=5903 auid=4294967295 uid=0 gid=0  
euid=0 suid=0 fsuid=0 egid=0 sgid=0 fsgid=0 tty=(none) ses=4294967295  
comm="runc:[2:INIT]" exe="/" key=(null)
```



# Soundness Challenge: In-depth Analysis of Namespace Implication

Namespace	Virtualized System Resource	What low-level events are affected?	Impact on soundness
PID	Process IDs	Events with PIDs	Yes, fragmentation
Mount	Mount points	Events with file pathnames	Yes, ambiguities
Network	Network device/stack, etc.	Events with Local IPs/ports	Yes, both
IPC	SYSV IPC objects & POSIX msg queue	Events with ID/names of SYSV IPC object/POSIX msg queues	Yes, ambiguities
User	User/group IDs	Events with GIDs and UIDs	No
Time, UTS, Cgroup	Boot/ monotonic clocks; Host/NIS domain name; Cgroup root directory	Does not affect provenance dataflow	N/A

PID/Mount/Network/IPC ns can impact soundness. Additional details provided in paper (Tables 1 and 2).

# Soundness Challenge is not specific to Linux Audit

Question: *Does soundness challenge persist on other Linux auditing subsystems/OS platforms?*

Namespace	Sysdig	LTTng
PID	Yes	No
Mount	Yes	Yes
Network	Yes	Yes
IPC	Yes	Yes


Namespace	BSD Jail	Solaris Zone
PID	No	No
Mount	Yes	Yes
Network	Yes	Yes
IPC	Yes	Yes

- Soundness challenge persists in other Linux auditing subsystems and OS platforms.
- LTTng provides virtualized PIDs and host PIDs together so fragmentation can be solved by correlation.
- No fragmentation is caused by PID in BSD Jail/Solaris Zone because all PIDs are host PIDs.
- Additional details provided in paper (Table 3 and 4).

# Soundness Challenge: Mapping Virtualized Namespaces

---

- *Key design:* establish a mapping between the host view and the container view on virtualized resources.
- The mapping correlates the virtualized data with the host data.
- It helps provenance tracking system to use the **consistent** provenance data.
- The principal namespaces that impact provenance are PID, Mount, Network, and IPC.




# Clarity:

## Essential Container Semantic Patterns

---

- Understanding essential container semantics helps by:
  1. accelerating the automatic analysis process.
  2. simplifying security analysis.



# Effectiveness: Real-world Exploit

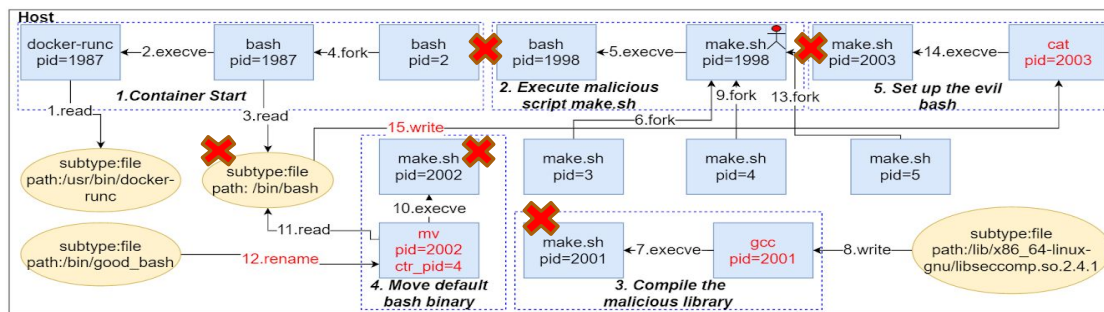
---

- Question: *How effective is CLARION in dealing with real-world container microservice attacks?*
- Answer: Real-world exploit

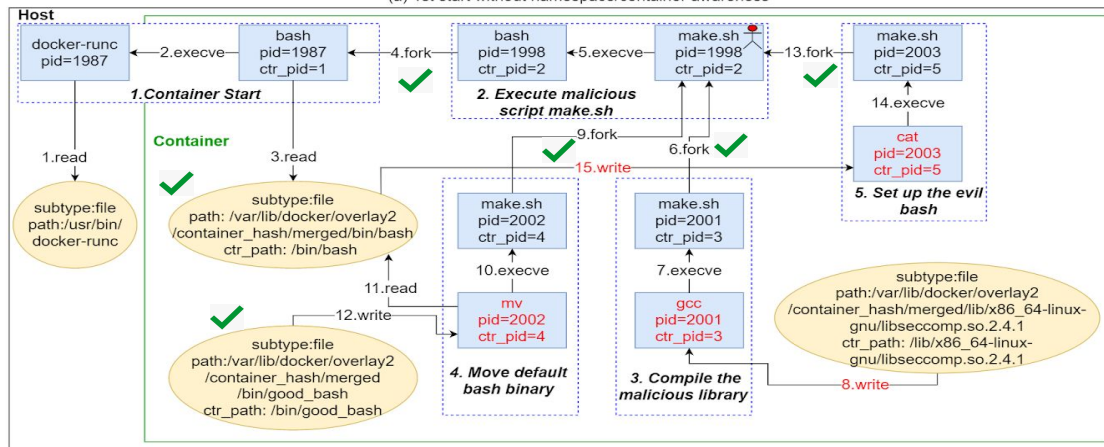
CVE #	Description	Severity
2019-5736	runC related	High. Achieve privilege escalation.
2019-14271	docker-tar related	High. Achieve privilege escalation.
2018-15664	docker-cp related	Normal. Achieve container host file system modification.

- We will focus on CVE 2019-5736 (runC) to show the effectiveness:
  - The exploit involves 2 starts of the compromised container.
  - The first start does the configuration.
  - The second start does the actual damage.

# Effectiveness: CVE 2019-5736

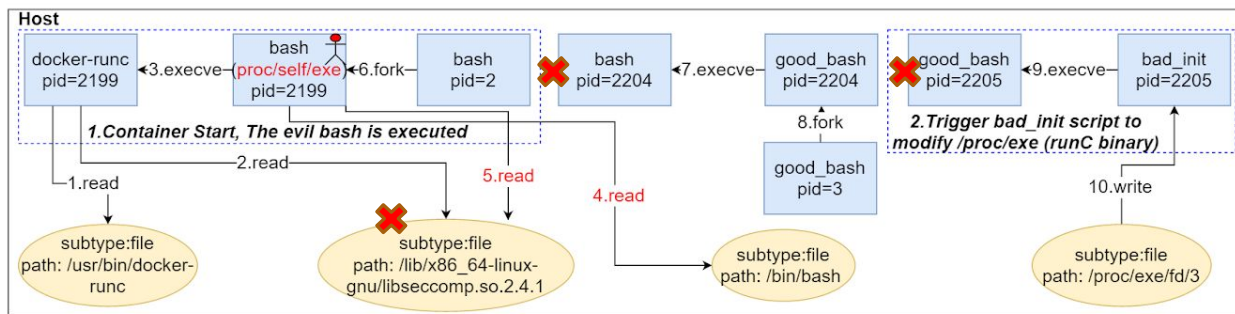


(a) 1st start without namespace/container awareness

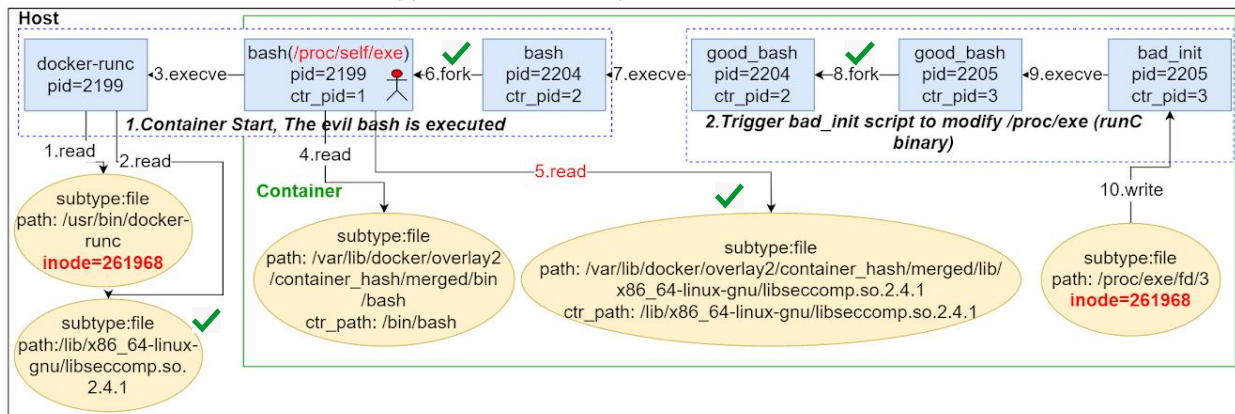


(b) 1st start with namespace/container awareness

# Effectiveness: CVE 2019-5736



(a) 2nd start without namespace/container awareness



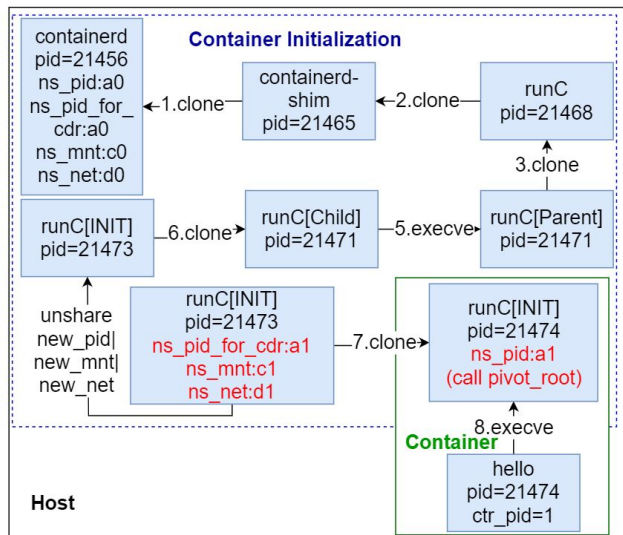
(b) 2nd start with namespace/container awareness

# Generality: Container Initialization

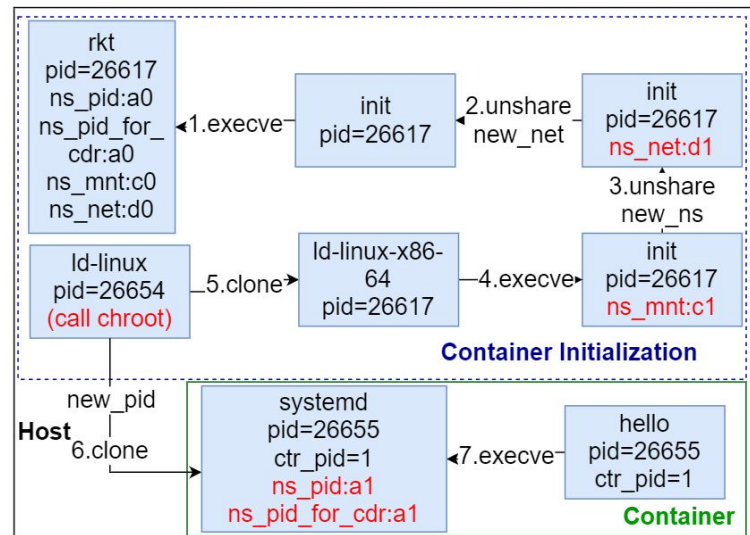
Question: *Is CLARION solution generally applicable across different container engines?*

Answer: Container initialization graph generation and quantitative results for different containers

## Docker Hello-world Init



## Rkt Hello-world Init





# Efficiency: Runtime/Storage Overhead

## Runtime Overhead Comparison of Container Provenance Systems

Service	Base (secs)	Linux Audit (secs)	SPADE (secs)	CLARION (secs)	Incremental Overhead (CLARION)	Overall Overhead (Audit + SPADE + CLARION)	Overall Overhead (CamFlow)	Overall Overhead (SEL-Audit)
frontend	1503 s	1550 s	1558 s	1578 s	1.3%	3.7%	4.8%	32.4%
productcatalog service	668 s	679 s	681 s	691 s	1.5%	3.4%	9.1%	25.0%
currencyservice	1104 s	1139 s	1153 s	1169 s	1.4%	5.9%	12.9%	8.5%
paymentservice	1082 s	1123 s	1126 s	1143 s	1.5%	5.6%	11.5%	9.7%
shippingservice	434 s	446 s	449 s	451 s	0.4%	3.9%	22.5%	25.8%
emailservice	929 s	960 s	1028 s	1068 s	3.9%	15.0%	1.2%	17.6%
checkoutservice	682 s	719 s	714 s	734 s	2.8%	7.6%	3.2%	13.9%
recommendation service	8726 s	9418 s	9337 s	9729 s	4.2%	11.5%	9.5%	19.5%
adservice	4438 s	4454 s	4518 s	4571 s	1.2%	3.0%	5.3%	8.5%
loadgenerator	200 s	208 s	212 s	215 s	1.4%	7.5%	20.4%	29.4%

## Storage Overhead Comparison

SEL-Audit	CamFlow	SPADE	CLARION	Incremental Overhead
168.79 GB	312.56 GB	174.68 GB	181.75 GB	4.05%

**Runtime/storage incremental overheads < 5%**

# Conclusion

---

- Existing provenance tracking solutions are inadequate for microservice scenarios
  - Namespace unawareness causes fragmentation and ambiguities (soundness).
  - Container unawareness leads to missing essential container semantics (clarity).
- CLARION
  - The first namespace-aware and container-aware provenance tracking solution
  - Comprehensive evaluation shows effectiveness, generality and efficiency of our solution
    - **Effectiveness:** We generate sound and clear provenance graphs of 3 real-world CVEs, which outperform the graphs generated by the traditional solution.
    - **Generality:** We show that our solution is independent of container engines by providing the container initialization graphs and quantitative results for 3 container engines.
    - **Efficiency:** We use a microservice benchmark to test the runtime/storage overhead of CLARION and find that the overhead is <5% over SPADE.